

# **CMP 1203**

## LECTURE 8

# RAID : Redundant Arrays of Independent Disks


- Early disk drives large and costly
- Required highly controlled environments – heat – damaged circuitry, humidity- built up static
- Head crashes/ other failures could damage entire disks
- Needed more reliable disks
- Patterson, Katz and Gibson : improve reliability and performance using a number of “inexpensive” small disks ; Redundant Arrays of Inexpensive Disks
- Inexpensive is a relative term : “Independent”



IBM RAMAC


# RAID Level 0 (RAID-0)

- Data blocks placed in stripes across several disk surfaces
- One record occupies several sectors across many disk surfaces
- Also called drive spanning/ block interleave data striping or disk striping
- No redundancy → best performance BUT very unreliable in case of faults
- Very inexpensive
- Very unreliable: as number of disks increases, probability of fault increases
- Recommended for non-critical data or data that doesn't change frequently and is backed up regularly, requires high-speed read/write and low cost e.g. video/image editing




# RAID -1

- Also disk mirroring
- Best failure protection
- Data is written on two sets of drives
- Second drive is mirror set/shadow set
- Acceptable performance
- Slower write performance than RAID-0 because data is written twice
- Faster read performance because can read from disk arm which is closest to target sector
- Good for transactions and applications that need high fault tolerance e.g. accounting, payroll
- Expensive because need twice as many disks to store given amount of information




# RAID -2

- Uses extreme data striping
- Writes one bit per strip (instead of blocks)
- Additional drives used for error correction using hamming codes
- Failed drive can be reconstructed from hamming drive and vice versa
- Acts as if it was one big drive since one data bit written per drive
- Need accurate synchronization else data becomes scrambled
- Hamming code generation is time consuming hence RAID-2 is too slow for commercial implementation
- Theoretical



# RAID-3


- Similar to RAID-2 but uses only one drive to hold a parity bit
- Same duplication and synchronization as RAID-2 but more economical because only 1 drive for protection



| Letter | ASCII     | Parity(even) |            |
|--------|-----------|--------------|------------|
|        |           | High Nibble  | Low Nibble |
| W      | 0101 0111 | 0            | 1          |
| E      | 0100 0101 | 1            | 0          |
| A      | 0100 0001 | 1            | 1          |
| T      | 0101 0100 | 0            | 1          |
| N      | 0100 1000 | 1            | 1          |
| E      | 0100 0101 | 1            | 0          |
| R      | 0101 0010 | 0            | 1          |

# RAID-4


- Also theoretical level like RAID-2
- RAID-0 with parity
- Parity introduces performance bottle neck : need to write parity bit as well
- So cant service multiple writes concurrently e.g. write to strip 3, 4,1 because you need to update parity block



PARITY 1-4 = (Strip 1) XOR (Strip 2) XOR (Strip 3) XOR (Strip 4)

# RAID-5

- RAID- 4 but spread parity disks across entire array
- Can service concurrent requests : provides best read throughput
- E.g. write to drive 4 strip 6 and drive 1 strip 7 because need different disk arms for both data and parity
- Most complex disk controller
- Best protection of all levels for least cost: commercially successful
- Used for file and application servers, email servers, database and web servers etc.




$$\text{PARITY 1-3} = (\text{Strip 1}) \text{ XOR } (\text{Strip 2}) \text{ XOR } (\text{Strip 3})$$



# RAID-6

- Most raid systems can tolerate at most one disk failure
- BUT most disk failures occur in clusters. **Why?**
- Disks manufactured at same time have same end of life
- Disk failures: events that affect all drives at same time e.g. power surge
- RAID-1 can handle multiple disk failures if both drive and mirror not wiped out
- RAID-6 uses two sets of error correction strips for every horizontal row of drives :One is parity, second uses reed-solomon code
- This increases storage cost




# Homework

- RAID DP (double parity)
- Hybrid RAID systems

# Optical Memory

- CD-ROM
- Polycarbonate disks covered with reflective aluminum film, then sealed with protective coating
- Compact discs written from center to outside edge using spiral track of bumps
- Bumps are called pits
- Spaced between pits called lands
- How do they work? <https://www.youtube.com/watch?v=H-jxTzFrnpg>



# **Self-Monitoring, Analysis and Reporting Technology (S.M.A.R.T)**

- Monitoring system for computer HDDs to detect and report various indicators of reliability to predict failures
- It's an interface between BIOS and storage device
- If enabled, BIOS can process information from the storage device and send warning messages about potential failures
- Avails self-test or maintenance routines

# Self-Monitoring, Analysis and Reporting Technology (S.M.A.R.T)

- Attributes to be monitored are set by manufacturer

## SMART Attributes

Examples of SMART attributes, showing the

| SMART attributes |                                   |           |       |       |        |          |              |
|------------------|-----------------------------------|-----------|-------|-------|--------|----------|--------------|
| ID               | Attribute name                    | Threshold | Value | Worst | Status | Raw Data | Raw Hex      |
| 01               | Raw Read Error Rate               | 51        | 100   | 100   | OK     | 0        | 000000000000 |
| 03               | Spin Up Time                      | 25        | 100   | 100   | OK     | 3136     | 000000000C40 |
| 04               | Start/Stop Count                  | 0         | 99    | 99    | OK     | 1653     | 000000000675 |
| 05               | Reallocated Sector Count          | 11        | 100   | 100   | OK     | 0        | 000000000000 |
| 07               | Seek Error Rate                   | 51        | 100   | 100   | OK     | 0        | 000000000000 |
| 08               | Seek Time Performance             | 15        | 100   | 100   | OK     | 0        | 000000000000 |
| 09               | Power-On Hours Count              | 0         | 100   | 100   | OK     | 316913   | 00000004D5F1 |
| 0A               | Spin-up Retry Count               | 51        | 100   | 100   | OK     | 0        | 000000000000 |
| 0C               | Power Cycle Count                 | 0         | 100   | 100   | OK     | 747      | 0000000002E8 |
| BF               | G-Sense Error Rate                | 0         | 93    | 93    | OK     | 79827    | 0000000137D3 |
| C2               | HDA Temperature                   | 0         | 124   | 82    | OK     | 38       | 000000000026 |
| C3               | Hardware ECC Recovered            | 0         | 100   | 100   | OK     | 5508     | 000000001584 |
| C4               | Reallocated Event Count           | 0         | 100   | 100   | OK     | 0        | 000000000000 |
| C5               | Current Pending Sector Count      | 0         | 100   | 100   | OK     | 0        | 000000000000 |
| C6               | Off-line Scan Uncorrectable Count | 0         | 100   | 100   | OK     | 0        | 000000000000 |
| C7               | UltraDMA CRC Error Rate           | 0         | 200   | 200   | OK     | 17       | 000000000011 |
| C8               | Write Error Rate                  | 51        | 100   | 100   | OK     | 0        | 000000000000 |
| C9               | Soft Read Error Rate              | 0         | 100   | 100   | OK     | 0        | 000000000000 |
| DF               | Load/Unload Retry Count           | 0         | 100   | 100   | OK     | 2        | 000000000002 |
| E1               | Load/Unload Cycle Count           | 0         | 70    | 70    | OK     | 305164   | 00000004A80C |
| FF               |                                   | 51        | 100   | 100   | OK     | 0        | 000000000000 |

SMART reported good status.

Source: sandisk.com

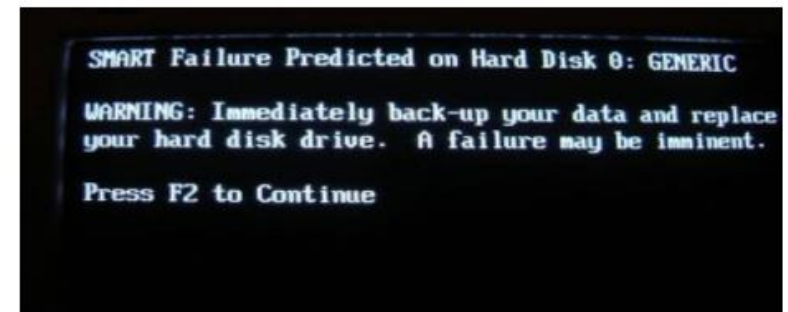
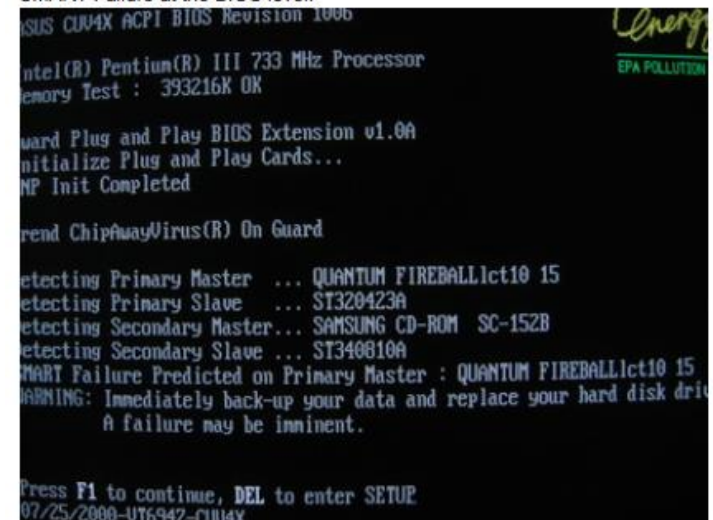
# Self-Monitoring, Analysis and Reporting Technology (S.M.A.R.T)

SMART failure at the Operating System level:



SMART Failure

SMART Failure at the BIOS level:



Source: sandisk.com

# Error Detection and Correction

- No channel or storage medium is 100% error free
- The higher the transmission rate, higher bit timing
- The more bits stored per square mm of storage, magnetic flux densities increase
- Error rate increases proportional to transmission rate and number of bits per sq. mm of storage
- We cant create an error free medium and we cant detect 100% of all possible errors
- Trade-off: define an acceptable number of errors for system to work. As long as we can detect this and correct this “reasonable” number, all is ok.
- “reasonable” differs per application/ implementation

# Error Detection: Parity method

- Uses an additional bit added to the code group → parity bit: either 0 or 1
- Even parity: choose bit so that the total number of 1's in the code group is even


e.g. **1 1 0 0 0 0 1 1** or **0 1 0 0 0 0 0 1**

- Odd parity : number of 1's is odd
- Can detect only single bit errors

**Qn:** Which is in error assuming odd parity?

**1 1 0 0 0 0 0 1** or **1 1 0 0 0 0 0 0**

**Qn:** Which bit was in error?





# Error Detection: Parity method

Qn: What if 2 bits are in error?

Ans: Doesn't change even or odd count of 1's so it won't identify that an error has occurred.

So when is it used?

- When probability of a single error is very low and that of double errors is 0. For example? (Homework!)
- Transmitter and receiver have to agree beforehand whether to use odd or even parity

# Error Detection: CRC

- Modulo 2 addition:  $0+0=0$ ;  $0+1=1$ ;  $1+0=1$ ;  $1+1=0$ .
- Modulo 2 division:
  - Write divisor directly beneath 1<sup>st</sup> bit of dividend
  - Add them using modulo 2
  - Bring down bits from dividend until 1<sup>st</sup> one of difference aligns with first 1 of divisor
  - Repeat until you reach number not divisible by divisor

$$\begin{array}{r}
 \text{quotient} \quad \mathbf{1010} \\
 \mathbf{1011} \overline{) 1001011} \\
 \underline{+ 1011 \downarrow} \\
 \mathbf{0100} \\
 \underline{+ 0000 \downarrow} \\
 \mathbf{1001} \\
 \underline{+ 1011 \downarrow} \\
 \mathbf{0101} \\
 \underline{+ 0000} \\
 \mathbf{101} \\
 \text{remainder}
 \end{array}$$

# Error Detection: CRC

- Operations have polynomial equivalents

$$1011_2 = 1 \times 2^3 + 0 \times 2^2 + 1 \times 2^1 + 1 \times 2^0$$

$$\text{Let } X = 2; \quad 1 \times X^3 + 0 \times X^2 + 1 \times X^1 + 1 \times X^0$$

- CRC uses such as generator polynomials
  - Let information be  $I = 1001011_2$
  - Sender and receiver agree on a binary pattern e.g.

$$P = 1011_2$$

# Error Detection: CRC

3. Shift I to left by one less than number of bits in P i.e.

$$I = 1001011000_2$$

4. Do modulo 2 division (work this out!) **Remainder??** Ans= $100_2$  becomes CRC check sum.
5. Add remainder to I to give message M [ $1001011100_2$ ]
6. At receiver, message is decoded by  $M \div P$ . If remainder is 0, no error. If remainder is another value, then there was an error.

# Example


- Key: 1101
- Received word: 10000001
- Determine if there is an error.
- What word will be sent if the message is 100100?

# Solution

Receiver Side

Let there be error in transmission media

Code word received at the receiver side - 100000001




Since the remainder is not all zeroes, the error is detected at the receiver side.

Data word to be sent - 100100

Key - 1101

Sender Side:



Therefore, the remainder is 001 and hence the code word sent is 100100001.

# Hamming Code

- In data communication systems, error detection is enough.
- If an error is detected, simply ask sender to re transmit.
- This is impossible with storage systems and memory.
- We must recover the data.
- Hamming codes are an efficient way
- Based on Parity
- Use check bits/ redundant bits
- Code word has  $n$  bits;  $n = m + r$  :  $m$  data bits and  $r$  check bits

# Hamming Code

- Define hamming distance as number of bit positions in which 2 code words differ.
- E.g. code words
- 10001001 and
- 10110001 have hamming distance of 3
- Hamming distance determines how many bit errors can be detected
- If distance is  $d$ , if  $d$  bits are flipped, error cannot be detected because it produces a valid code word
- Smallest hamming distance  $D(\min)$  is the smallest distance between all pairs of code words
- To detect  $k$  errors, need minimum distance of  $k+1$
- To correct  $k$  errors, minimum hamming distance must be at least  $2k+1$



# Example

| Data Word | Parity Bit | Code word |
|-----------|------------|-----------|
| 00        | 0          | 000       |
| 01        | 1          | 011       |
| 10        | 1          | 101       |
| 11        | 0          | 110       |

But 3 bits can have 8 possible combinations;  
only above 4 are valid words

|     |     |  |
|-----|-----|--|
| 000 | 100 |  |
| 001 | 101 |  |
| 010 | 110 |  |
| 011 | 111 |  |

- Assume even parity
- If we receive 111; invalid code word hence error has occurred
- But we cant correct it
- Cant tell how many bits flipped and which in error
- What if 2 bits flipped? A valid code word is generated:  $D(\min)$  is 2: So can detect only single bit errors

# Example 2

- Consider code below
- $D(\min) = 3$  (compare all code words)
- Can detect 2 errors and correct 1 bit error
- What if we receive 10000?
- Find hamming distance between each code word [1,4,2,3]
- Choose legal code word closest to received word
- It may not necessarily be correct. We have assumed 1 bit error occurred
- What if we received 11000 and we know 2 bits flipped? [2,3,3,2]
- We don't know closest code word!

|           |
|-----------|
| 0 0 0 0 0 |
| 0 1 0 1 1 |
| 1 0 1 1 0 |
| 1 1 1 0 1 |

# Hamming Code Generation

- Use inequality:  $(m + r + 1) \leq 2^r$
- Specifies lower limit of how many check bits are needed. (For single bit errors)
- Homework: How do we derive this inequality?
- If data words have  $m = 4$  bits, then

$$(4 + r + 1) \leq 2^r \quad ; r \geq 3$$

- So we need 3 check bits to build a code word to correct single bit errors for a 4 bit data word

# Hamming Code Generation

1. Determine number of check bits  $r$  required.
2. Number the  $n$  bits from right to left starting with 1
3. Each bit whose position is a power of 2 is a parity bit : rest are data bits
4. Parity bits check specific data bits such that: bit  $b$  is checked by the parity bits  $b_1, b_2, \dots, b_j$  such that  $b_1 + b_2 + \dots + b_j = b$  (  $+$  is modulo 2 sum)

# Example

- Encode data word 01001011 using a hamming code
- Step 1: Determine code word
- $m = 8; \quad (8 + r + 1) \leq 2^r \quad ; r \geq 4$
- Step 2: number bits from right to left

|    |    |    |   |   |   |   |   |   |   |   |   |
|----|----|----|---|---|---|---|---|---|---|---|---|
|    |    |    |   | X |   |   |   | X |   | X | X |
| 12 | 11 | 10 | 9 | 8 | 7 | 6 | 5 | 4 | 3 | 2 | 1 |

- X are parity bit positions (power of 2 positions)

# Hamming Code Generation

- Step 3: Assign parity bits to check bit positions. Write bit positions as sums of numbers which are powers of 2

|       |         |          |
|-------|---------|----------|
| 1=1   | 5=1+4   | 9=1+8    |
| 2=2   | 6=2+4   | 10=2+8   |
| 3=1+2 | 7=1+2+4 | 11=1+2+8 |
| 4=4   | 8=8     | 12=4+8   |

- 1 appears in 1,3,5,7,9,11 so it will be show parity for those positions
- 2 will check 2,3,6,7,10,11 etc

# Hamming Code Generation

- Write data words in blank positions and fill parity
- Assume even parity

01001011

|    |    |    |   |   |   |   |   |   |   |   |   |
|----|----|----|---|---|---|---|---|---|---|---|---|
| 0  | 1  | 0  | 0 | X | 1 | 0 | 1 | X | 1 | X | 0 |
| 12 | 11 | 10 | 9 | 8 | 7 | 6 | 5 | 4 | 3 | 2 | 1 |

- 1: 1,3,5,7,9,11
- 2: 2,3,6,7,10,11
- 4: 4,5,6,7,12
- 8: 8,9,10,11,12

Fill in the rest

# Hamming Code Generation

- Write data words in blank positions and fill parity
- Assume even parity

|    |    |    |   |   |   |   |   |   |   |   |   |
|----|----|----|---|---|---|---|---|---|---|---|---|
| 0  | 1  | 0  | 0 | 1 | 1 | 0 | 1 | 0 | 1 | 1 | 0 |
| 12 | 11 | 10 | 9 | 8 | 7 | 6 | 5 | 4 | 3 | 2 | 1 |

- 1: 1,3,5,7,9,11
  - 2: 2,3,6,7,10,11
  - 4: 4,5,6,7,12
  - 8: 8,9,10,11,12
- Code word for K: 010011010110



# Hamming Code Generation

- Assume error in position 9: 010**1**11010110- received
- Even parity 010**0**11010110 - codeword

|    |    |    |   |   |   |   |   |   |   |   |   |
|----|----|----|---|---|---|---|---|---|---|---|---|
| 0  | 1  | 0  | 1 | 1 | 1 | 0 | 1 | 0 | 1 | 1 | 0 |
| 12 | 11 | 10 | 9 | 8 | 7 | 6 | 5 | 4 | 3 | 2 | 1 |

- 1: 1,3,5,7,9,11 : error
- 2:2,3,6,7,10,11 : ok
- 4:4,5,6,7,12 :ok      Code word for K: 010011010110
- 8: 8,9,10,11,12 :error

# Hamming Code Generation

- Assume error in position 9: 010111010110

|    |    |    |   |   |   |   |   |   |   |   |   |
|----|----|----|---|---|---|---|---|---|---|---|---|
| 0  | 1  | 0  | 1 | 1 | 1 | 0 | 1 | 0 | 1 | 1 | 0 |
| 12 | 11 | 10 | 9 | 8 | 7 | 6 | 5 | 4 | 3 | 2 | 1 |

• 1: 1,3,5,7,9,11 : error

2: 2,3,6,7,10,11 : ok

• 4: 4,5,6,7,12 :ok

8: 8,9,10,11,12 :error

• Parity bits 1 and 8 have errors: common bits are 9 and 11

• But 11 checked by parity 2 and is ok

• So error must be on 11

• We can also just sum parity bits in error  $(1+8) = 9$  to get bit position with error

# Home work

- Reed Solomon codes

# Handout

- Lobur and Null; Chapter 2 and 7
- Stallings, Chapter 6 , 5
- NOTE: CAT 2 - 15<sup>th</sup> April